

Analysis of Forking Processes in a UNIX Environment

Abraham J. Reines

February 14, 2024

Abstract

This report explains the process of executing a program intended to demonstrate the forking of processes in a UNIX environment. Running a C program to create multiple processes using the forking command and system calls facilitated an understanding of differences between on-screen and redirect output to files with and without the use of `fflush(stdout)`. The findings demonstrate the technicalities of process management in concurrent systems. This draws attention to the use of buffer management for output behavior.

1 Introduction

The process of forking is fundamental in UNIX and POSIX-compliant operating systems. This report aims to explore the behavior of the `fork` system call and its implications on process management and output buffering.

2 Methodology

The C program is compiled and executed on a UNIX system. The output is observed on-screen, redirected to a file, and appended to another file to analyze the behavior of process output under different conditions.

3 Code Explanation

The provided C program¹ utilizes the `fork()` system call, which is fundamental for creating child processes in UNIX. When `fork()` is called, the current process is duplicated, creating a child process with a new process ID. The child process receives a copy of the parent's memory, but as execution continues, their states may diverge. The `getpid()` and `getppid()` functions retrieve the process ID and parent process ID, respectively, which are used in the output to track process lineage. The use of `fflush(stdout)` ensures that the standard output buffer is flushed immediately, affecting the order and visibility of the output across different processes.

¹See next page for the full code listing.

3.1 With fflush

The use of fflush(stdout) ensures that the output buffer is flushed after each print statement, which affects how often the output is written to the terminal or file. *See next page for code.

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 int main(void)
5 {
6     int k;
7     printf ("Main Process' PID = %d\n", getpid());
8     fflush(stdout);
9     for (k = 1; k <= 3; k++)
10    {
11        fork ();
12        printf ("k = %d, PPID = %d, pID = %d, I'm Alive!\n", k, getppid(), getpid());
13        fflush(stdout);
14    }
15 }
```

3.2 Without fflush

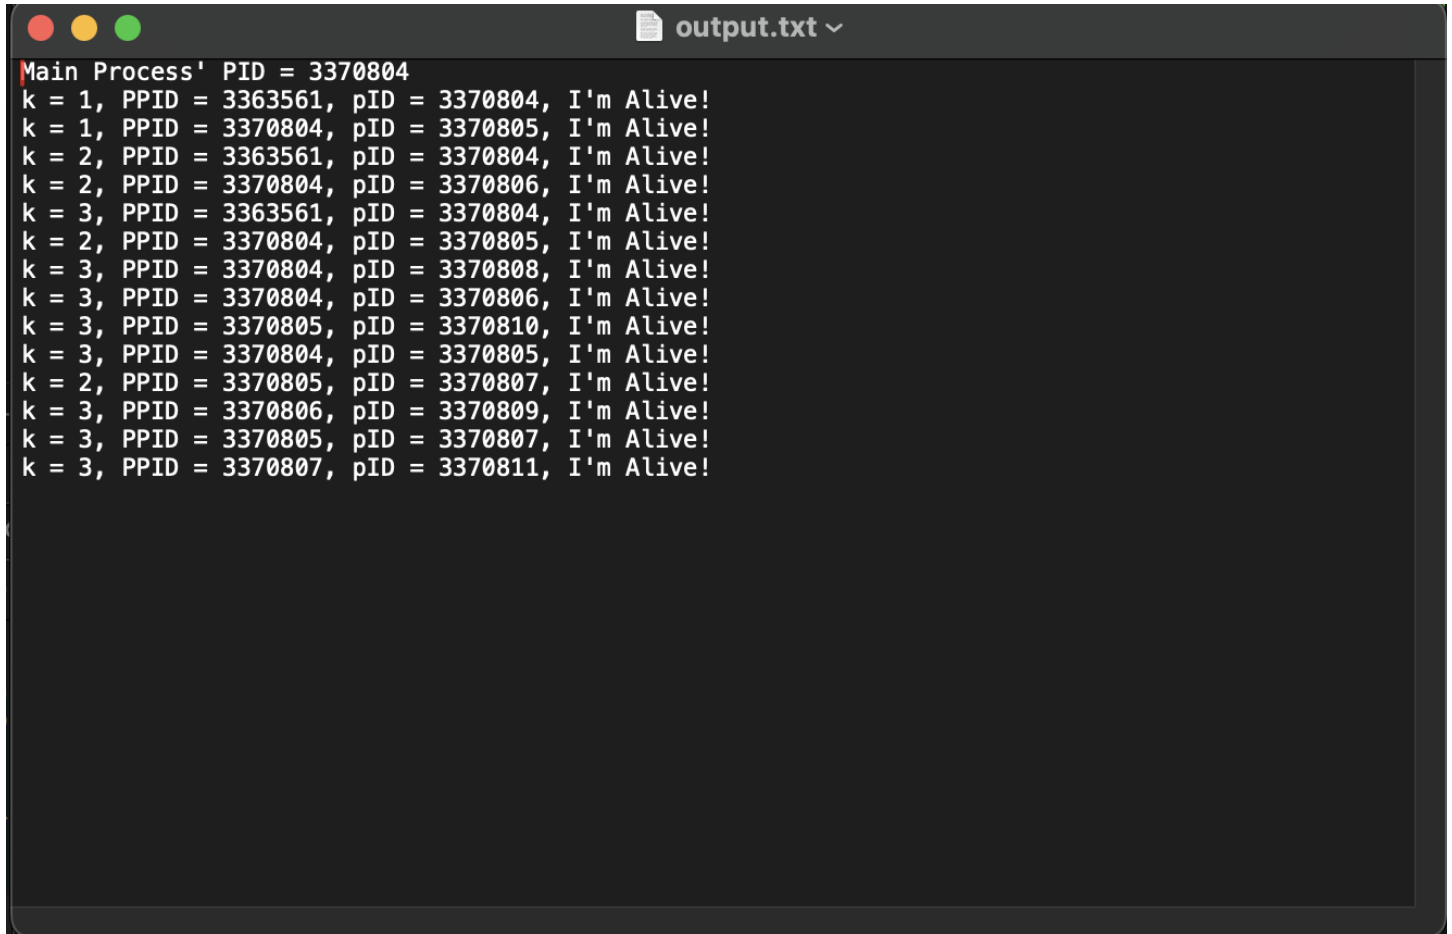
Commenting out the fflush(stdout) statement changes the program's behavior due to the output buffer not being manually flushed after each print statement.

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 int main(void)
5 {
6     int k;
7     printf ("Main Process' PID = %d\n", getpid());
8     fflush(stdout);
9     for (k = 1; k <= 3; k++)
10    {
11        fork ();
12        printf ("k = %d, PPID = %d, pID = %d, I'm Alive!\n", k, getppid(), getpid());
13        //fflush(stdout);
14    }
15 }
```

4 Results and Analysis

This section will include the screenshots of the terminal output and the contents of the output files to illustrate the effects of forking and output buffering.

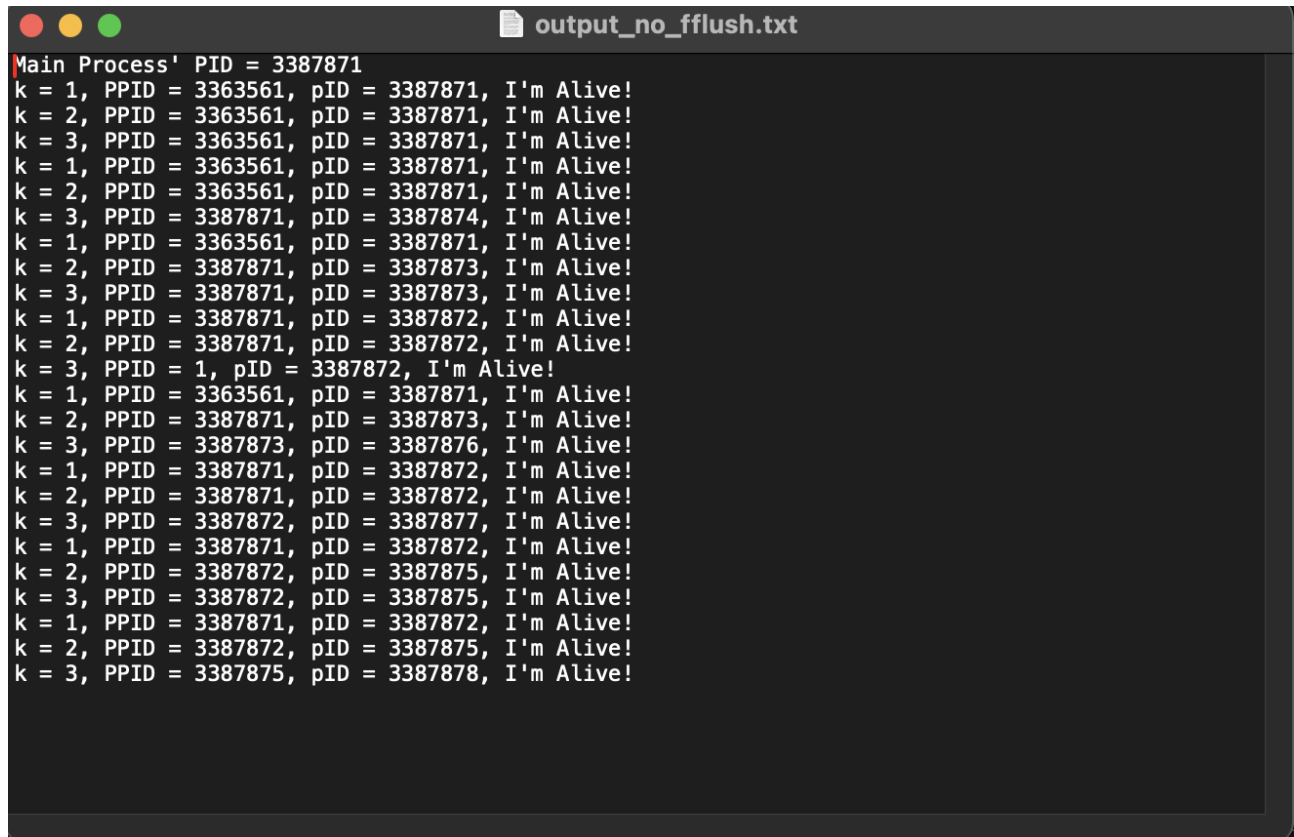
4.1 Terminal Output with fflush

A screenshot of a terminal window with a dark background. The title bar at the top shows three colored window control buttons (red, yellow, green) on the left and a file icon followed by the text 'output.txt' and a dropdown arrow on the right. The terminal content shows the output of a program. The first line is 'Main Process' PID = 3370804'. This is followed by 15 lines of output, each starting with 'k =', followed by 'PPID =', 'pID =', and 'I'm Alive!'. The values for k, PPID, and pID vary across the lines, showing the state of different processes and their parents. The output is displayed line-by-line, indicating that fflush is being used to flush the output buffer immediately.

```
Main Process' PID = 3370804
k = 1, PPID = 3363561, pID = 3370804, I'm Alive!
k = 1, PPID = 3370804, pID = 3370805, I'm Alive!
k = 2, PPID = 3363561, pID = 3370804, I'm Alive!
k = 2, PPID = 3370804, pID = 3370806, I'm Alive!
k = 3, PPID = 3363561, pID = 3370804, I'm Alive!
k = 2, PPID = 3370804, pID = 3370805, I'm Alive!
k = 3, PPID = 3370804, pID = 3370808, I'm Alive!
k = 3, PPID = 3370804, pID = 3370806, I'm Alive!
k = 3, PPID = 3370805, pID = 3370810, I'm Alive!
k = 3, PPID = 3370804, pID = 3370805, I'm Alive!
k = 2, PPID = 3370805, pID = 3370807, I'm Alive!
k = 3, PPID = 3370806, pID = 3370809, I'm Alive!
k = 3, PPID = 3370805, pID = 3370807, I'm Alive!
k = 3, PPID = 3370807, pID = 3370811, I'm Alive!
```

Figure 1: Terminal output when fflush is used.

4.2 Terminal Output without fflush



```
output_no_fflush.txt
Main Process' PID = 3387871
k = 1, PPID = 3363561, pID = 3387871, I'm Alive!
k = 2, PPID = 3363561, pID = 3387871, I'm Alive!
k = 3, PPID = 3363561, pID = 3387871, I'm Alive!
k = 1, PPID = 3363561, pID = 3387871, I'm Alive!
k = 2, PPID = 3363561, pID = 3387871, I'm Alive!
k = 3, PPID = 3387871, pID = 3387874, I'm Alive!
k = 1, PPID = 3363561, pID = 3387871, I'm Alive!
k = 2, PPID = 3387871, pID = 3387873, I'm Alive!
k = 3, PPID = 3387871, pID = 3387873, I'm Alive!
k = 1, PPID = 3387871, pID = 3387872, I'm Alive!
k = 2, PPID = 3387871, pID = 3387872, I'm Alive!
k = 3, PPID = 1, pID = 3387872, I'm Alive!
k = 1, PPID = 3363561, pID = 3387871, I'm Alive!
k = 2, PPID = 3387871, pID = 3387873, I'm Alive!
k = 3, PPID = 3387873, pID = 3387876, I'm Alive!
k = 1, PPID = 3387871, pID = 3387872, I'm Alive!
k = 2, PPID = 3387871, pID = 3387872, I'm Alive!
k = 3, PPID = 3387872, pID = 3387877, I'm Alive!
k = 1, PPID = 3387871, pID = 3387872, I'm Alive!
k = 2, PPID = 3387872, pID = 3387875, I'm Alive!
k = 3, PPID = 3387872, pID = 3387875, I'm Alive!
k = 1, PPID = 3387871, pID = 3387872, I'm Alive!
k = 2, PPID = 3387872, pID = 3387875, I'm Alive!
k = 3, PPID = 3387875, pID = 3387878, I'm Alive!
```

Figure 2: Terminal output when fflush is not used.

4.3 File Output Analysis

In output buffering, this compares terminal and file outputs with and without the use of `fflush`. `fflush` results in immediate output to the terminal, while without `fflush`, the output is buffered, leading to a non-deterministic order in terminal output. Buffered I/O may lead to race conditions.

5 Conclusion

The experiment demonstrates the impact of the `fork` system call on process creation and management. It also highlights the role of output buffering in the visibility of print statements to the terminal and files.

Academic Integrity Pledge

“This work complies with the JMU honor code. I did not give or receive unauthorized help on this assignment.”